



Data Engineer Take Home Exam

General instructions

You are given **48 hours** upon receiving the exam to complete the deliverables.

Remember : **do not upload any parts of this exam or your answers on publicly accessible sites** like the public repositories on Github. This is to ensure that the exam will not be leaked to future applicants.

Upon completion, send an email to your hiring manager with a link to your deliverables to stop the timer. Good luck!

Part 1: Algorithmic Thinking

Euler discovered the remarkable quadratic formula:

$$n^2 + n + 41$$

It turns out that the formula will produce 40 primes for the consecutive integer values $0 \leq n \leq 39$. However, when $n = 40$, $40^2 + 40 + 41 = 40(40 + 1) + 41$ is divisible by 41, and certainly when $n = 41$, $41^2 + 41 + 41$ is clearly divisible by 41.

The incredible formula $n^2 - 79n + 1601$ was discovered, which produces 80 primes for the consecutive values $0 \leq n \leq 79$. The product of the coefficients, -79 and 1601 , is -126479 .

Considering quadratics of the form:

$$n^2 + an + b, \text{ where } |a| < 1000 \text{ and } |b| \leq 1000$$

where $|n|$ is the modulus/absolute value of n

e.g. $|11| = 11$ and $|-4| = 4$

Find the product of the coefficients, a and b , for the quadratic expression that produces the maximum number of primes for consecutive values of n , starting with $n = 0$.

Please code up a brute-force solution to this problem (it will not finish in a reasonable time). If you pass the exam stage, prepare to explain possible solutions to improve your current approach in relation to runtime during the panel interview.

Part 2: Soft(ware) skills

We live and breathe data. So naturally, we try to make our internal processes as data-driven as possible. Time-tracking allows us to determine our teams' bandwidth and throughput so we checkin religiously. We do this using [a handful of tools](#) that we can use without interrupting our work.

For this exam, you are given [an anonymized dump of our checkins](#) for the past year. Your goal is to:

- Clean the data
- Load it to your choice of database
- Create a web service that returns the checkin data associated to a given user

Additionally, answer the following questions:

- If the data is to be ingested periodically, what changes will you make to your current approach?
- Draw a data architecture showing different components of your ETL process.
- How will you verify the correctness of the ingested data?

Bonus points for:

- Using Python
- A publicly accessible deployment of your service
- Documentation
- Tests
- Diagrams
- Creating a web service that displays a per-user filtered view of the check-ins

We value:

- Communication
- Reproducibility
- Pragmatism
- Code hygiene

Submit your code by sharing a private Github repo with the following users:

- <https://github.com/blaise-tm>
- <https://github.com/tm-jamie-macagba>
- <https://github.com/tm-vincent-carpio>
- <https://github.com/tm-flo-barot>
- <https://github.com/gab-tm>
- <https://github.com/tm-kobe-deluna>
- <https://github.com/jillianchan>
- <https://github.com/tm-justin-beredo>

As a backup, please also upload your code to the Drive folder shared with you.